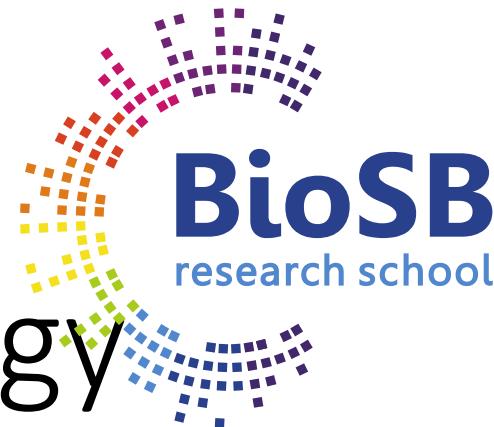


# Machine Learning for Bioinformatics & Systems Biology



## 5. Deep learning strategies & Generative Modelling

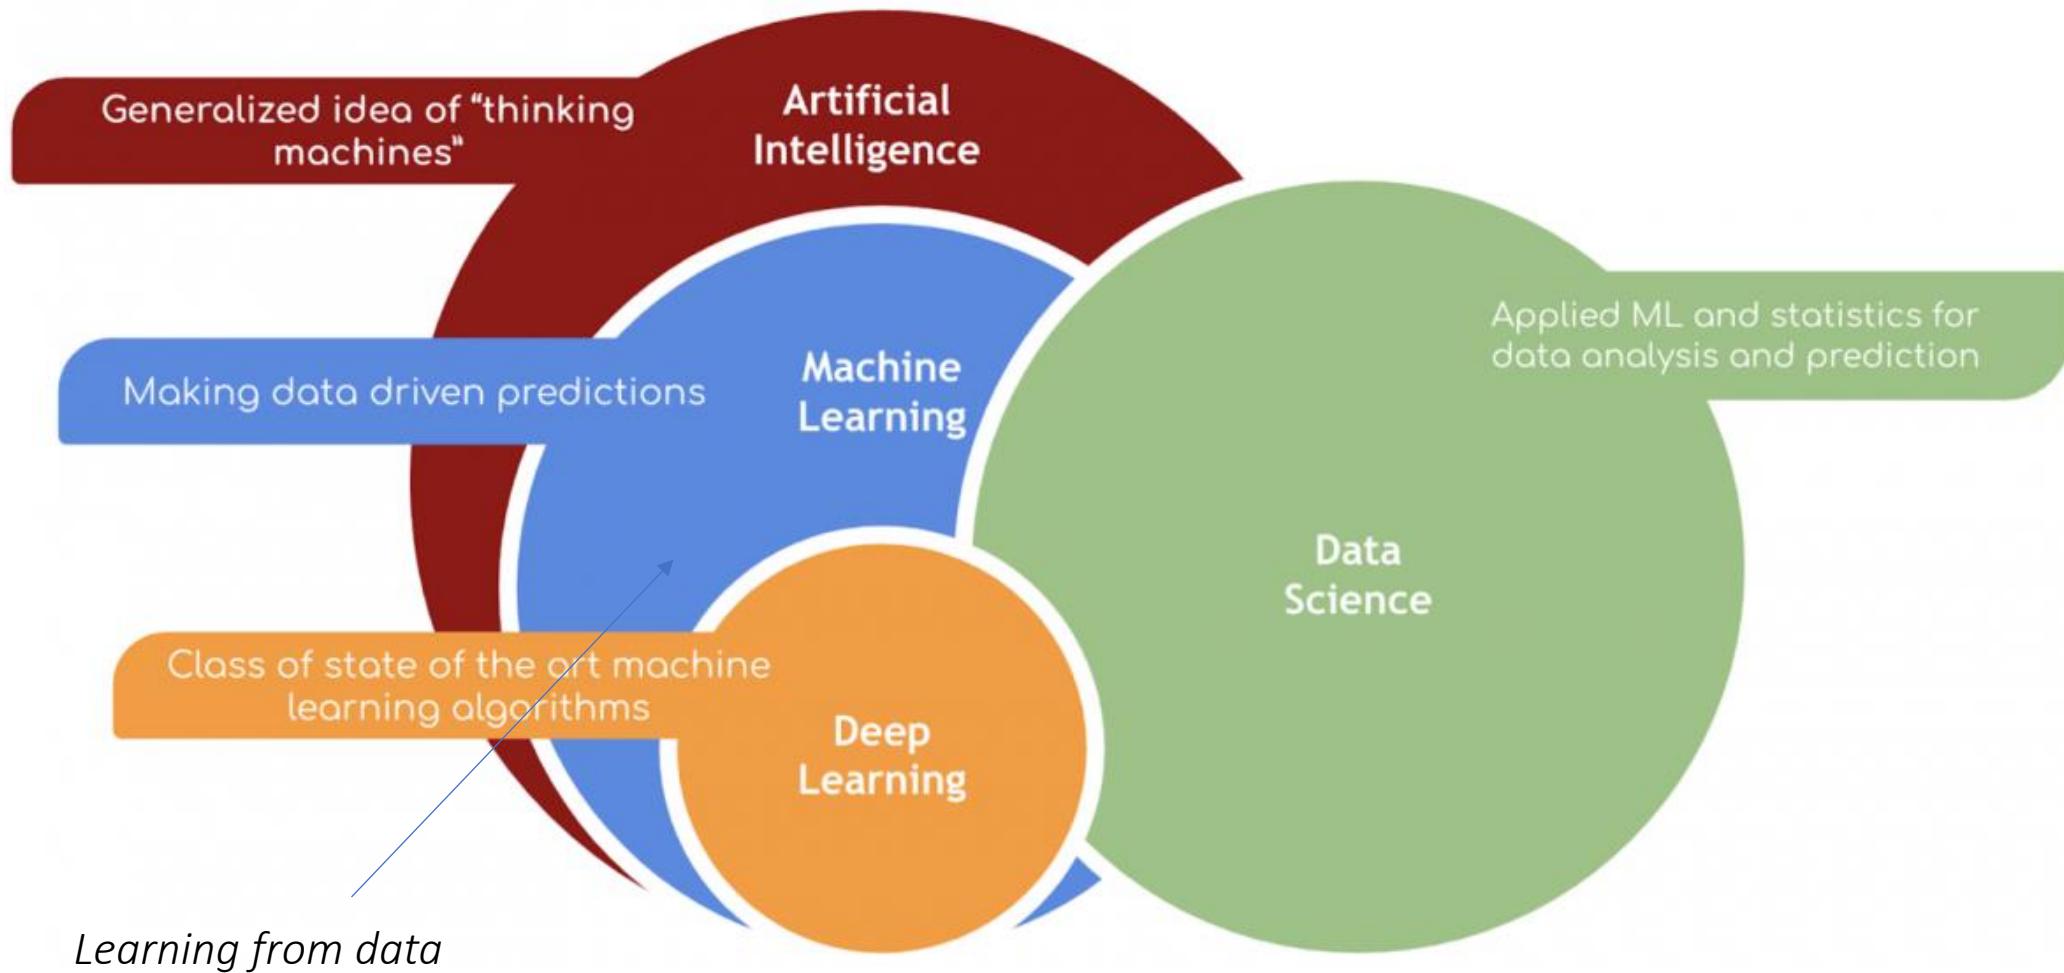
Marcel Reinders *Delft University of Technology*

Perry Moerland *Amsterdam UMC, University of Amsterdam*


Lodewyk Wessels *Netherlands Cancer Institute*

# Deep learning strategies

# What is Artificial Intelligence?



# Machine learning : learning from data



(height)

(level metabolite X)

Statistical modeling

$Y_i = a + bX_i + e_i$

Neural network modeling

$Y_i$



$b$

$a$

1

$X_i$

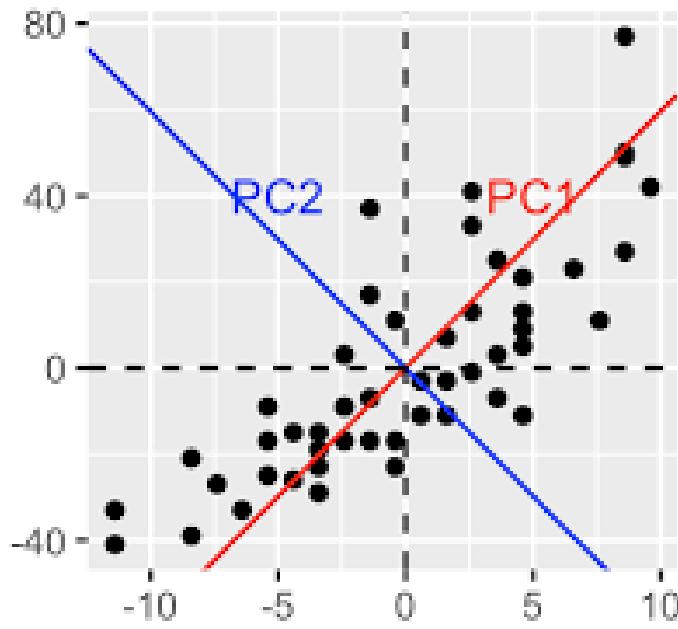
# Complex functions, non-linear complex networks




*But you need a lot of (labeled) data ( $Y$ 's)*

# Reduce dimensionality, autoencoder


## Build predictor in reduced space, *embeddings*



PCA  
linear reduction



Autoencoder  
non-linear reduction

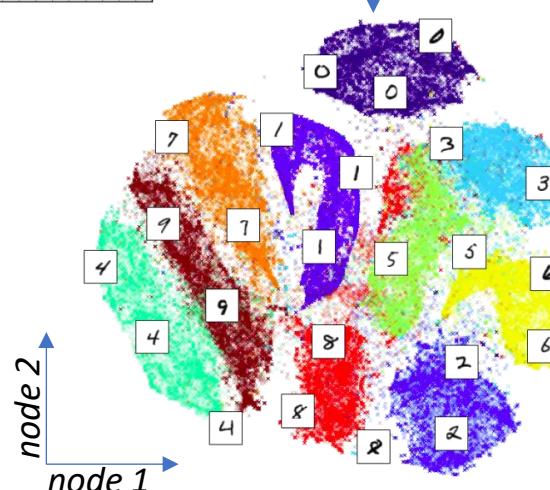
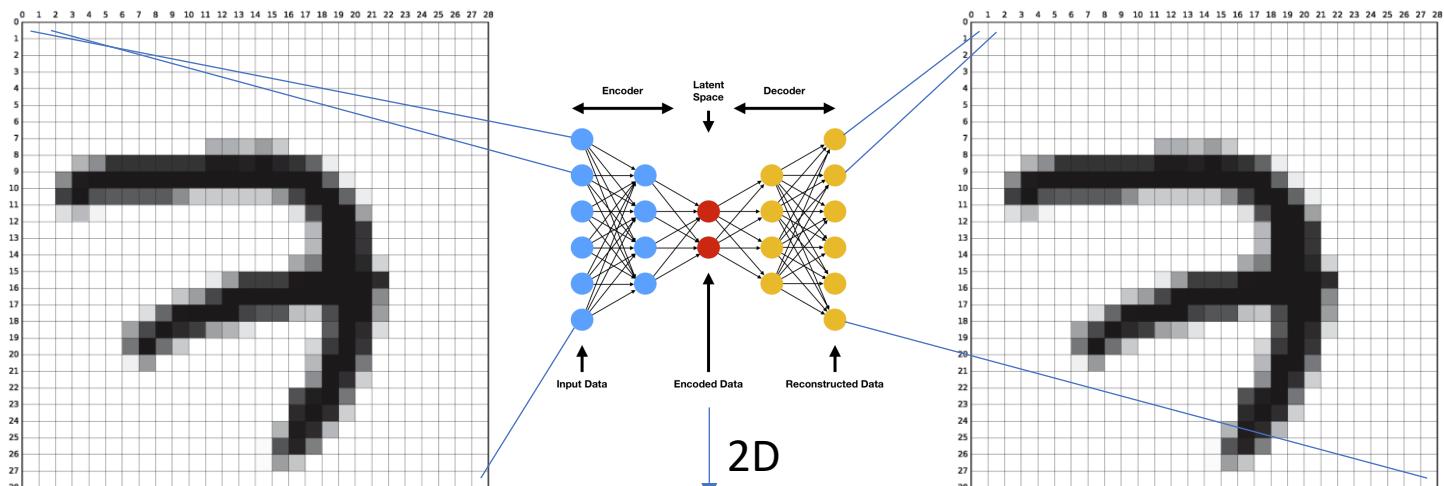


latent space, embedding, learned representation

Use learned representation to predict  $Y$



# Famous example: MNIST

0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9




Learned representation  
Digits separable

# Multi modalities: cross-modality imputation, joint spaces

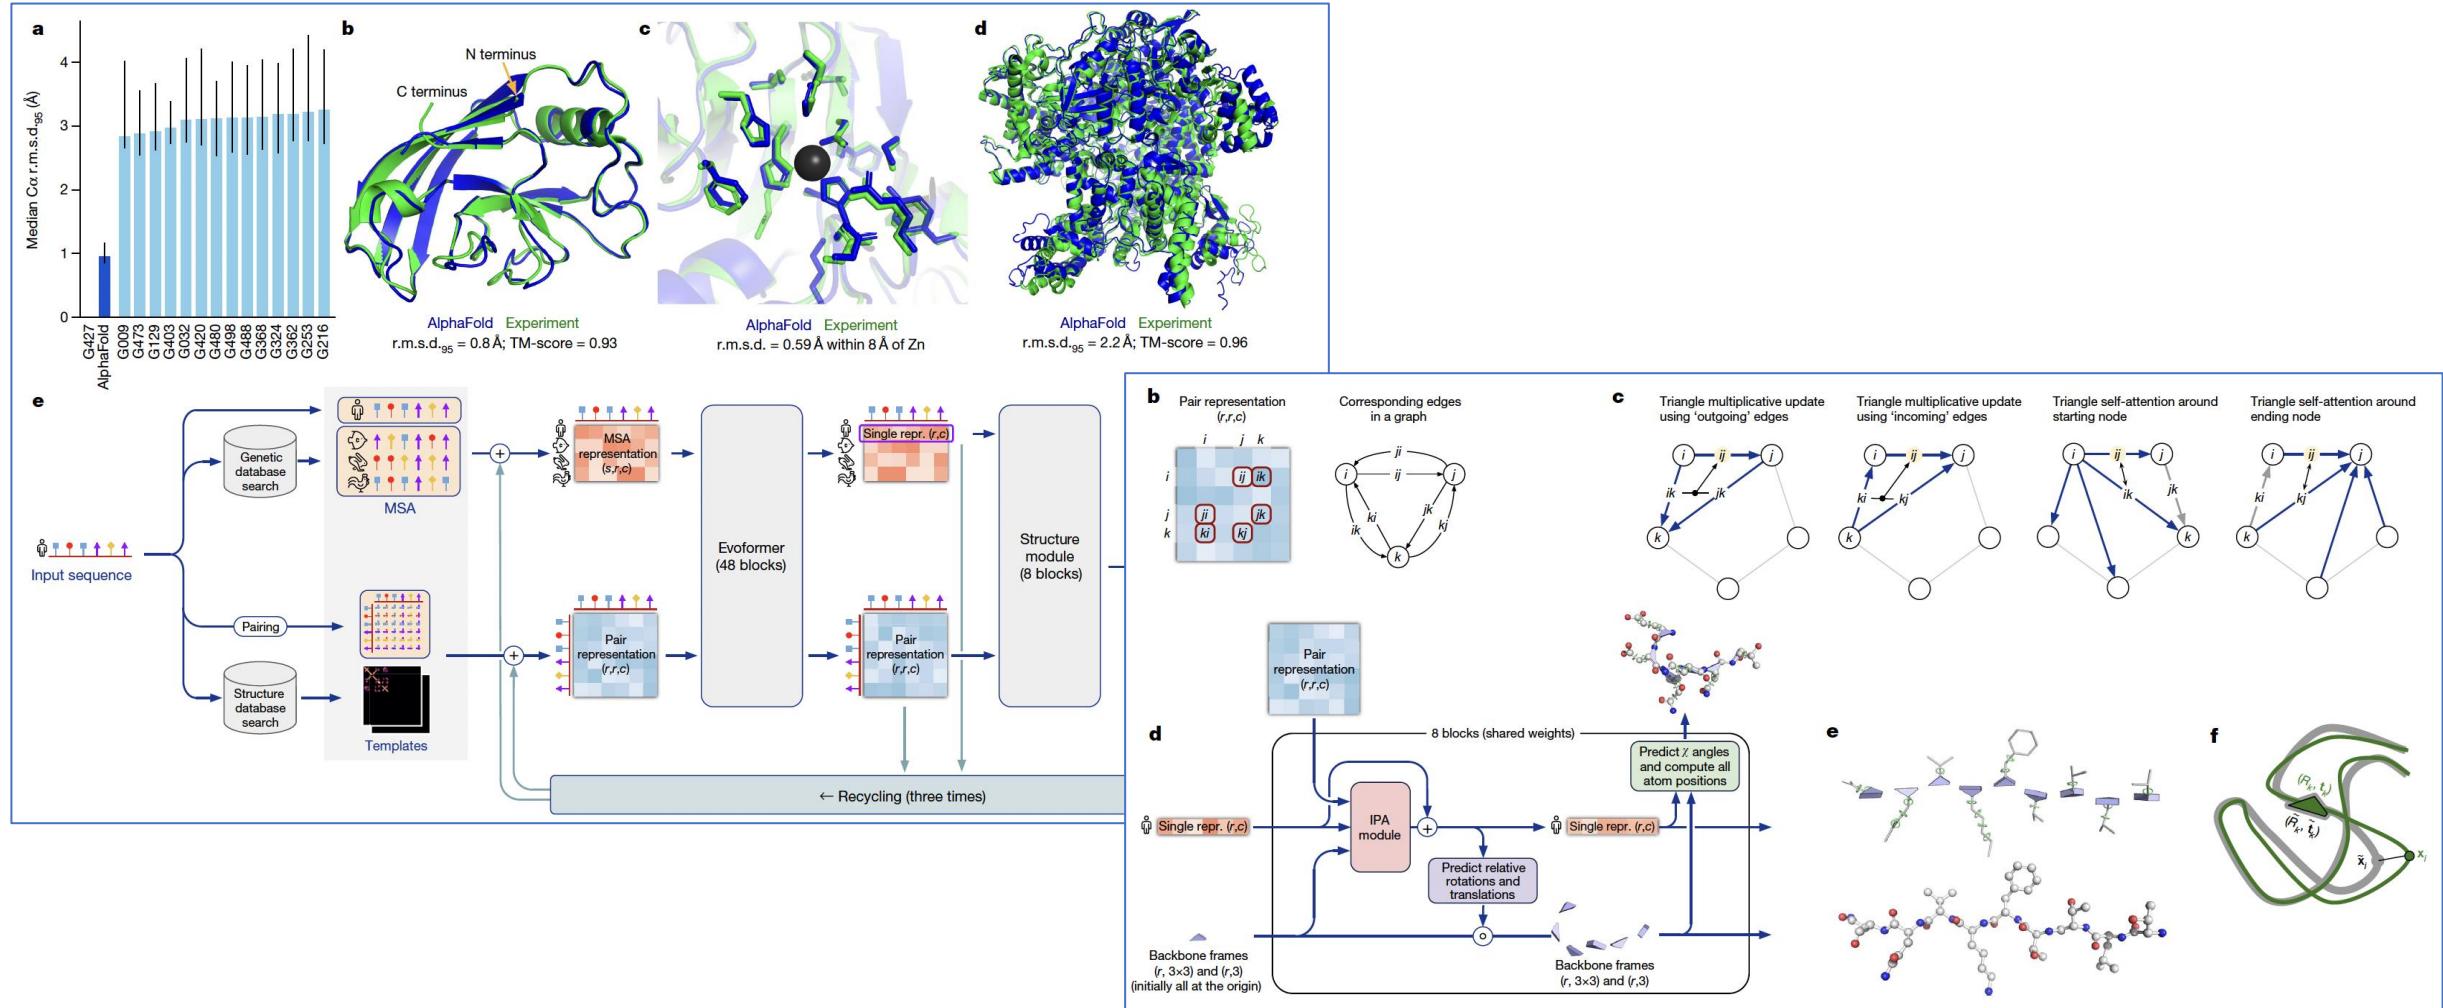


# Incorporate knowledge into neural network

Eg how molecules relate to each other




# AlphaFold: predicting 3D structure of proteins Based on graph convolutional neural net (GNN)




# Incorporate knowledge into neural network


*Eg how genes relate to each other*



**Incorporate knowledge** into neural network  
*Eg on physical constraints on output*

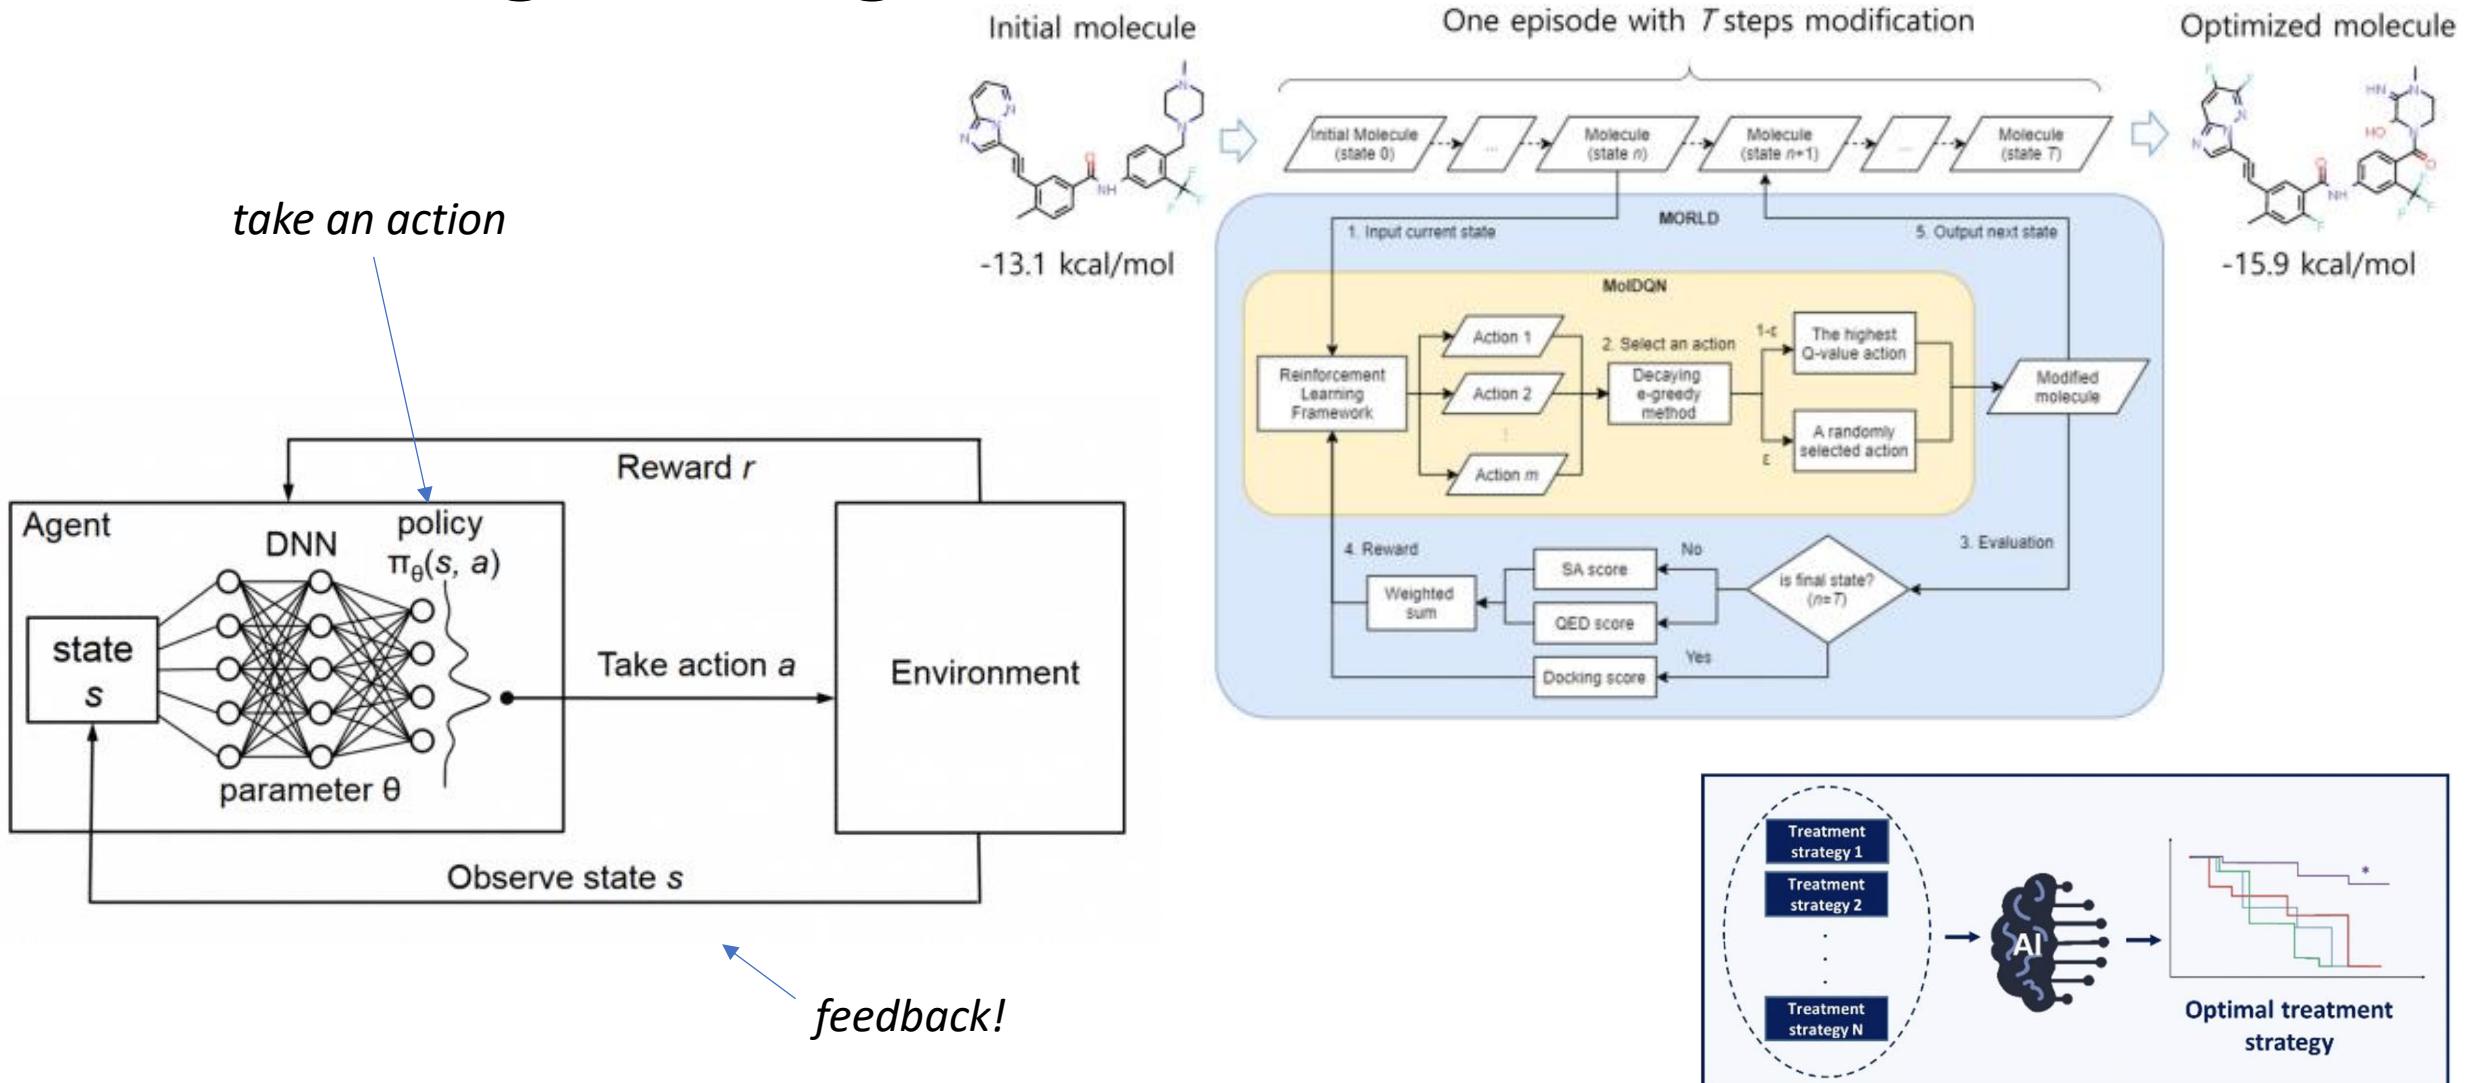


# Learn cellular drug/perturbation response



# Reinforcement learning

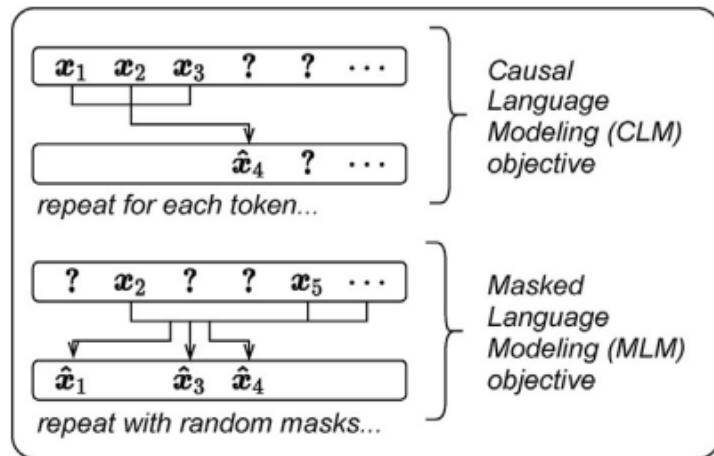
## Learning strategies



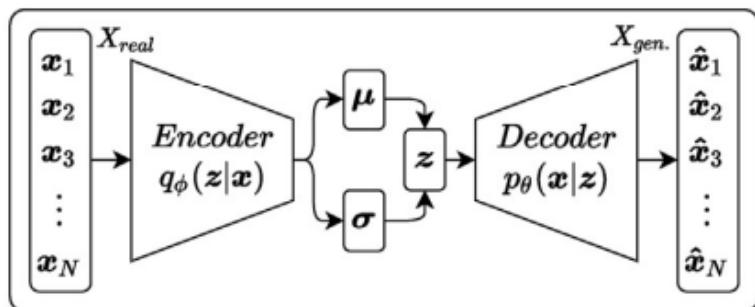
# Generative Modelling

# Deep generative model families

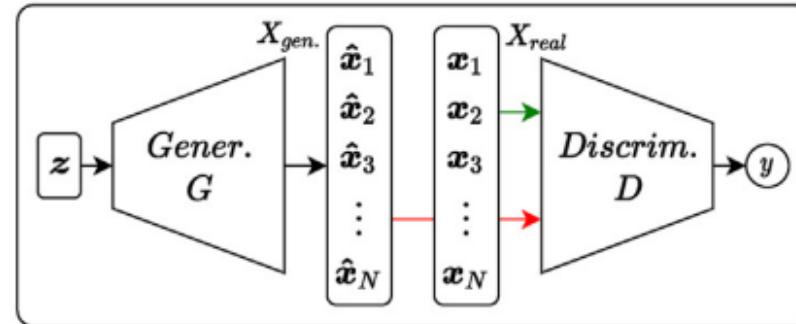
## Autoregressive language models (LMs)




## Variational AutoEncoders (VAEs)




## Generative Adversarial Networks (GANs)



## Normalising Flows (NFs)



## Diffusion Models (DMs)



# Applications in Bioinformatics

## Autoregressive language models (LMs)

e.g., protein design (ProtGPT2, ProGen, ...), drug discovery (molGPT, ...), DNABert, ...

## Variational AutoEncoders (VAEs)

e.g., protein design, drug discovery, representation learning in transcriptomics, multimodal data integration, ...

## Generative Adversarial Networks (GANs)

e.g., data generation/augmentation across different bioinformatics domains (scRNA-seq data, DNA sequences, Protein sequences), ...




## Diffusion Models (DMs)

e.g., protein folding, protein and small molecule generation, protein-ligand interaction modeling, data analysis (cryo-EM, gene expression), single-cell image, ... and many more...

## Normalising Flows (NFs)


# Autoregressive Language Model

# Autoencode sequences (Language models)



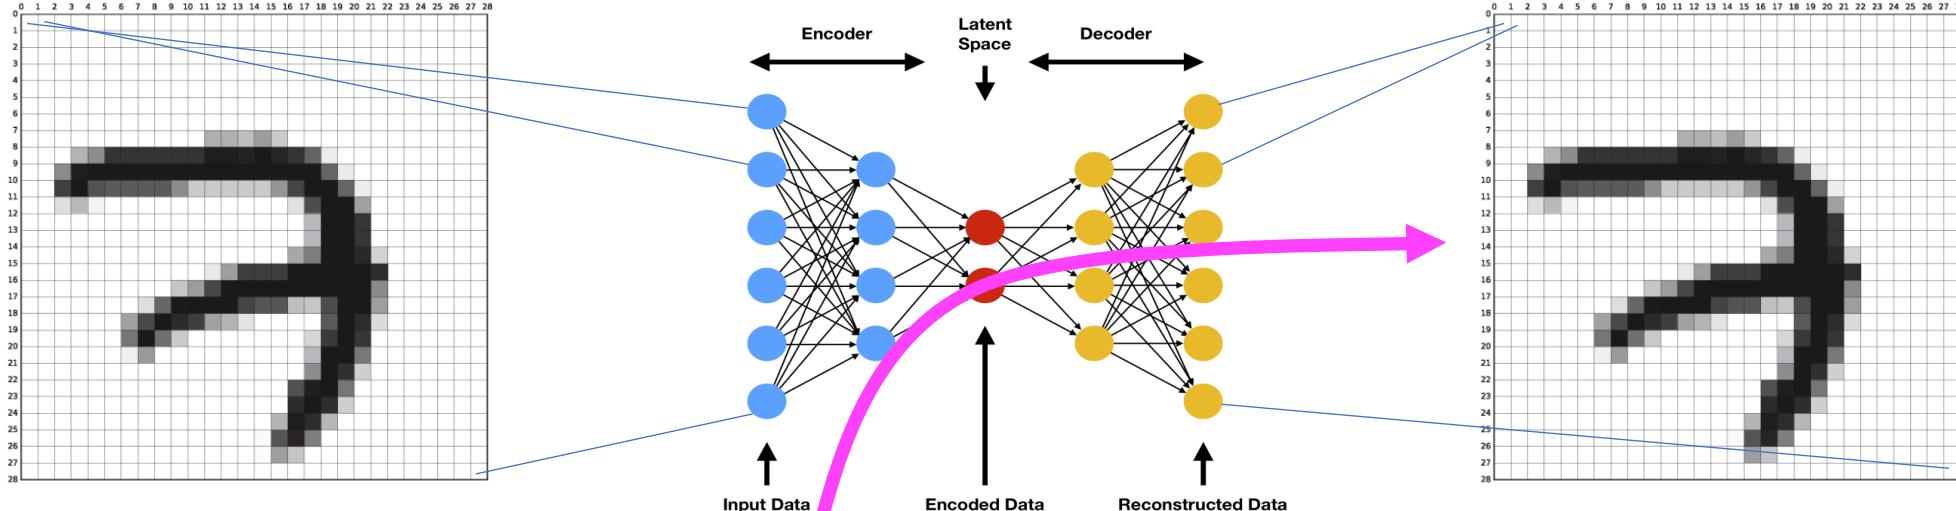
# Autoencode sequences (Language modes)

## Protein embeddings, function prediction, redesign



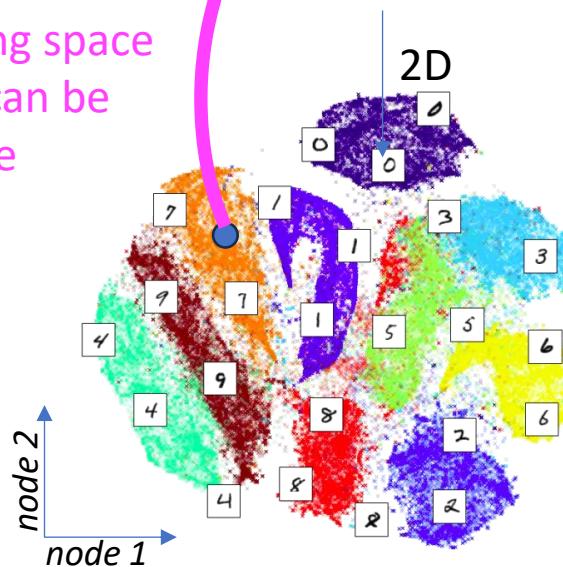
# Variational Autoencoder

# Revisit Autoencoder



Sampling from embedding space generates samples that can be decoded and look like the samples in input space

Enforce this behavior more:  
Require that samples close to  
embedded sample should  
generate the same sample




Learned representation  
Digits separable

# Probabilistic encoder


Choice prior can be anything, Gaussian, Laplacian, Student t,  
Mixture of gaussians, etc.

$$z \sim \mathcal{N}(\mu(x), \sigma^2(x)) \quad \equiv \quad z = \mu(x) + \sigma(x)\epsilon, \epsilon \sim \mathcal{N}(0, 1)$$



# Backpropagation problem

- Backpropagation cannot flow through the random node
- Solve by reparameterization



Instead of sampling from  
 $z \sim q_\phi(z | x)$

Sample from  
 $\epsilon \sim N(0, 1)$

And linear transform  
 $z = \mu + \sigma \odot \epsilon$

Alternatively, sampling from  
 $z \sim N(\mu, \sigma)$

Is the same as sampling from  
 $\epsilon \sim N(0, 1)$

And setting

$$z = \mu + \sigma \odot \epsilon$$

# What to optimize

- For each sample  $\mathbf{z}$ , there will be two variables  $\mu$  and  $\sigma$  (*defining Gaussian*)
- Accumulation of all Gaussian distributions becomes the original distribution  $P(\mathbf{x})$

$$P(x) = \int_z P(z) P(x|z) dz$$

- $P(\mathbf{x})$ , likelihood of the data, the bigger the better

$$\text{Maximum } L = \sum_x \log P(x)$$

LogP(x)

$$\log P(x) = \int_z q(z|x) \log P(x) dz = \log P(x) \boxed{\int_z q(z|x) dz}$$

equals 1 (trick)

$$= \int_z q(z|x) \log \left( \frac{P(z, x)}{P(z|x)} \right) dz$$

$$= \int_z q(z|x) \log \left( \frac{P(z, x)}{q(z|x)} \frac{q(z|x)}{P(z|x)} \right) dz$$

$$= \int_z q(z|x) \log \left( \frac{P(z, x)}{q(z|x)} \right) dz + \boxed{\int_z q(z|x) \log \left( \frac{q(z|x)}{P(z|x)} \right) dz}$$

$$= \int_z q(z|x) \log \left( \frac{P(x|z)P(z)}{q(z|x)} \right) dz + \boxed{KL(q(z|x)||P(z|x))}$$

Kullback Leiber divergence, always bigger than zero

$$\log P(x) \geq \int_z q(z|x) \log \left( \frac{P(x|z)P(z)}{q(z|x)} \right) dz = \text{ELBO}$$

Maximizing  $P(x)$  equal to maximizing ELBO

$$\begin{aligned}
\text{ELBO} &= \int_z q(z|x) \log \left( \frac{P(x|z)P(z)}{q(z|x)} \right) dz \\
&= \int_z q(z|x) \log \left( \frac{P(z)}{q(z|x)} \right) dz + \int_z q(z|x) \log P(x|z) dz \\
&= \boxed{-KL(q(z|x)||P(z))} + \boxed{\int_z q(z|x) \log P(x|z) dz}
\end{aligned}$$

*maximize*      *Maximize*

$$\begin{aligned}
-KL(q(z|x)||P(z)) &= \int_z q(z|x) (\log P(z) - \log q(z|x)) dz \\
&= -\frac{1}{2} \sum_{i=1}^J (1 + \log(\sigma_i^2) - \mu_i^2 - \sigma_i^2)
\end{aligned}$$

$$\begin{aligned}
\int_z q(z|x) \log P(z) dz &= \int_z N(z; \mu, \sigma^2) \boxed{\log N(z; 0, I)} dz \\
\int_z q(z|x) \log q(z|x) dz &= \int_z N(z; \mu, \sigma^2) \log N(z; \mu, \sigma^2) dz
\end{aligned}$$

*prior on z*

$$\text{Maximum } \int_z q(z|x) \log P(x|z) dz$$

**P(x|z)** (Decoder's output) given that **q(z|x)** (Encoder's output) is as high as possible. Similar to AutoEncoder's loss function(reconstruction error):

$$(x - \hat{x})^2$$

# Some variations (many more)

- $\beta$ -VAE : Balance between reconstruction loss and KL term, learn to entangle ( $\beta>1$ )

$$L(\theta, \phi; x) = -\mathbb{E}_{q(z|x)}[\log p(x|z)] + \boxed{\beta D_{KL}(q(z|x) || p(z))}$$

- $\beta$ -Total Correlation VAE ( $\beta$ -TCVAE) : Additionally penalizing the total correlation between the latent variables, more statistically independent latent variables  



$$\text{mutual information between the data}$$

$$L(\theta, \phi; x, \alpha, \beta, \gamma) = -\mathbb{E}_{q(z|x)}[\log p(x|z)] + \alpha D_{KL}(q(z, x)||q(z)p(x)) + \beta D_{KL}(q(z)||\Pi q(z)) + \gamma \sum D_{KL}(q(z)||p(z))$$


*dependence term*
*mutual information between the data variable and latent variables*
  
*prevent latent variable to diverge from prior*

# Autoencoder with generating distribution

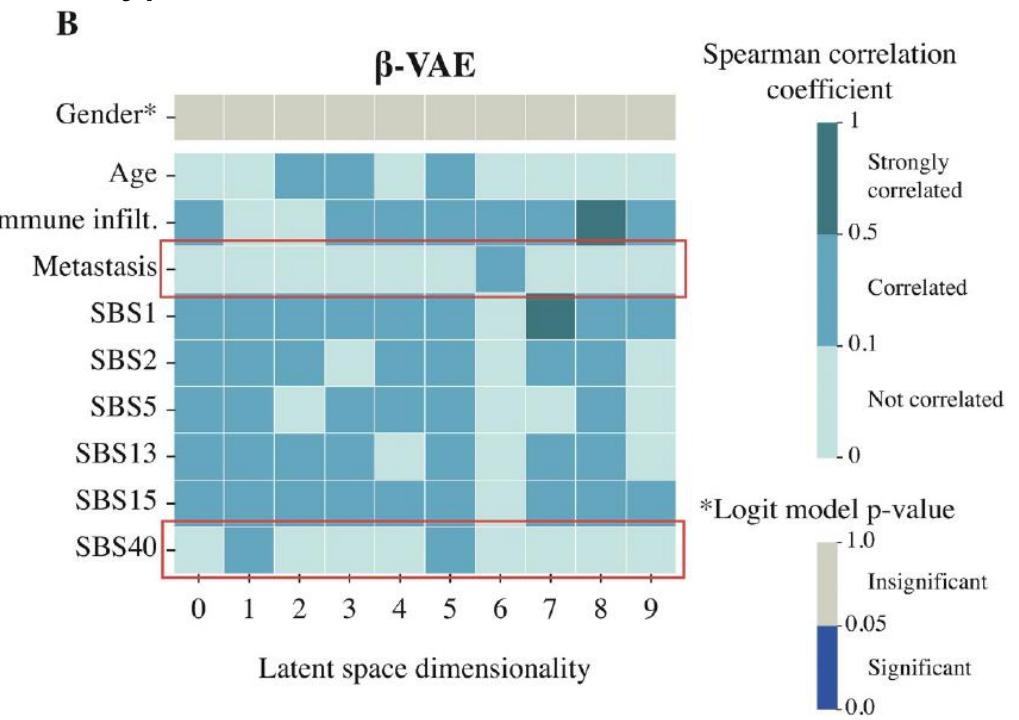
## Latent dimensions can vary different traits



# Example: Beta-VAE on GTEx data ; latent factors are disentangled and correlate with data features



17,382 samples with 56,200 genes  
representing 30 different tissue types



# Diffusion model

Idea: Recursively add noise (the diffusion)

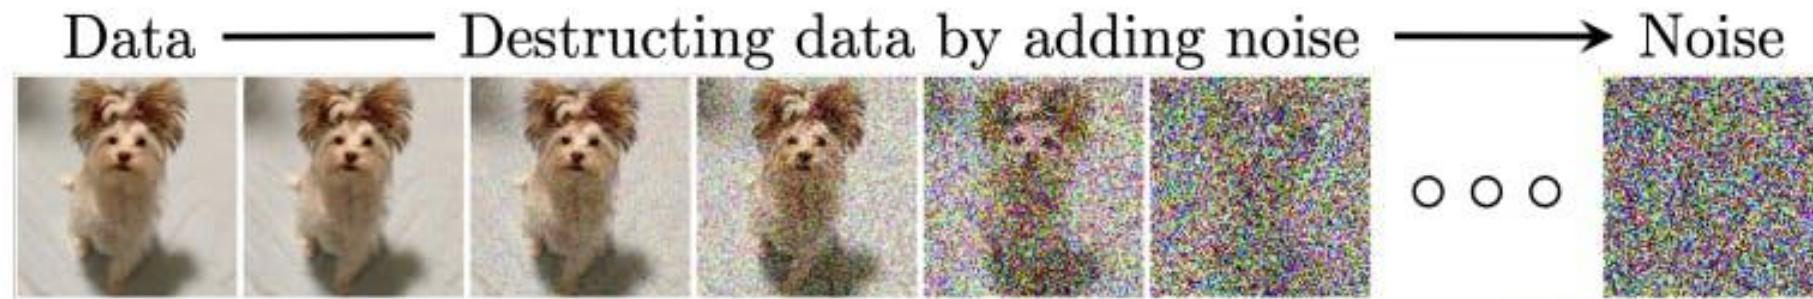
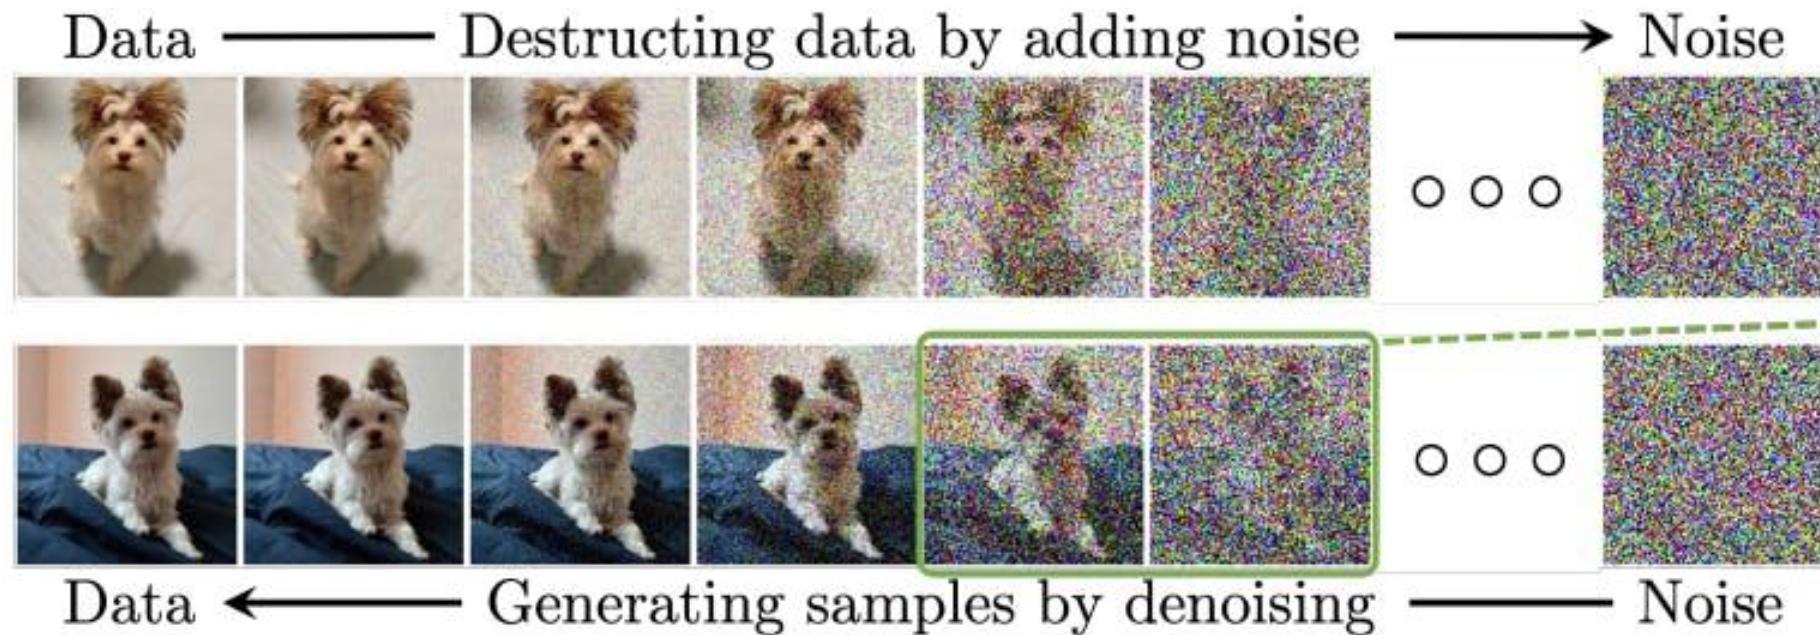


Image is **transformed** to TV static (white noise)

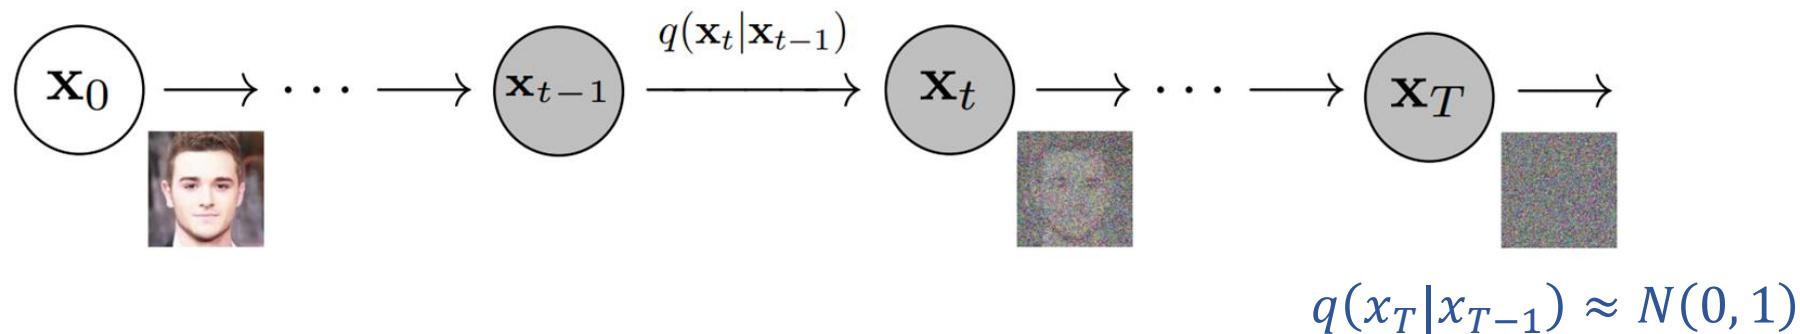
Idea: What if we could reverse this process?




Then we can draw random noisy image, and generate an image like our training data


# Diffusion Probabilistic Models

- **Forward diffusion process:** Markov chain gradually adds noise to data to obtain approximate posterior




- **Generative diffusion process:** use the Markovian assumption to learn the reverse process





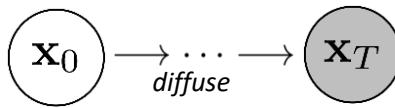
# Intuitive

- Need to learn  $q(\mathbf{x}_t | \mathbf{x}_{t-1})$  which is our model  $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$



- Provide any NN architecture with  $(\mathbf{X}_{t-1}, \mathbf{X}_t)$  and learn to predict  $\mathbf{x}_{t-1}$  from  $\mathbf{x}_t$ 
  - Want to optimize that for all diffusion steps
  - Need to do that for many diffusions, and for all time steps (to be able to generate from noise)

*more formal  
(sorry, some more math)*



# The forward process (diffusion process) (2)

The posterior distribution after T steps

$$q(\mathbf{x}_{1:T} | \mathbf{x}_0) := \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1}) := \prod_{t=1}^T \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$


The product of each state at time  $t$  given the previous state  $t-1$

Corrupt data by sampling from a multivariate Gaussian distribution mean centered around the previous state

Where  $\beta_1, \dots, \beta_T$  is a variance schedule (either learned or fixed)

Too large  $\beta$  will corrupt image too quickly. Very difficult to undo

Too small will take a long time to learn



# The forward process (diffusion process) (3)

- Property of the forward process is that it admits sampling  $x_t$  at an arbitrary timestep  $t$  in closed form (*one step instead of sequence of steps*)

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

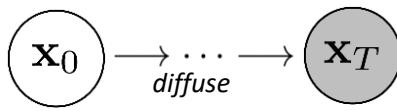
$$\alpha_t := 1 - \beta_t \quad \bar{\alpha}_t := \prod_{s=1}^t \bar{\alpha}_s$$

- Allows to rewrite, *reverse step* (!) (remember  $q$  was defined with:  $q(\mathbf{x}_t | \mathbf{x}_{t-1})$ )

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I})$$

$$\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0 + \frac{\sqrt{\bar{\alpha}_t} (1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t \quad \tilde{\beta}_t := \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$

Available, can be used for training!



# The reverse process (reverse diffusion process)

The parameterised Markov chain mapping noise back to image

$$p_{\theta}(\mathbf{x}_{0:T}) := p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) := p(\mathbf{x}_T) \prod_{t=1}^T \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}(\mathbf{x}_t, t), \Sigma_{\theta}(\mathbf{x}_t, t))$$

$$p(\mathbf{x}_T) = q(\mathbf{x}_T | \mathbf{x}_0) \approx \mathcal{N}(0, 1)$$


Different time steps are associated with different noise levels

$$\boxed{\mu_{\theta}(\mathbf{x}_t, t), \Sigma_{\theta}(\mathbf{x}_t, t)}$$

These need to be learned

Note: Reverse diffusion transition distribution depends only on the previous timestep (Markov):

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}(\mathbf{x}_t, t), \Sigma_{\theta}(\mathbf{x}_t, t))$$



# Reverse Markov transitions (1)

- Reverse Markov transitions:

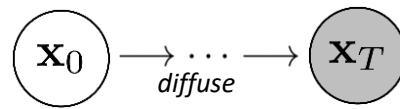
$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_{\theta}(\mathbf{x}_t, t), \boldsymbol{\Sigma}_{\theta}(\mathbf{x}_t, t))$$

- Find the one that maximize the likelihood of the training data

$$L_{vlb} = L_0 + L_1 + \dots + L_{T-1} + L_T$$

- Use Kullback-Leibler Distance to measures distance between two distributions

$$L_{t-1} = D_{KL}(q(x_{t-1}|x_t, x_0) \parallel p_{\theta}(x_{t-1}|x_t))$$



# Reverse Markov transitions (2)

- Simplify: Set the variances equal to variances in forward process schedule :

$$\Sigma_\theta(x_t, t) = \sigma_t^2 \mathbb{I} \quad \sigma_t^2 = \beta_t$$

- Reverse Markov transition becomes:

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t)) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \sigma_t^2 \mathbb{I})$$

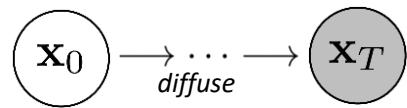
- Allows to transform (rewrite) Kullback-Leiber distance to:

$$L_{t-1} \propto \|\tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t)\|^2$$

known      model that we want

- The most straightforward parameterization of  $\mu_\theta(x_t, t)$  is a model that predicts the  $\tilde{\mu}_t(x_t, x_0)$  forward process posterior mean (known/training!)

$$\tilde{\mu}_t(x_t, x_0) := \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t$$



# Some rewriting

- Reparametrize forward step  $q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$

$$\mathbf{x}_t(\mathbf{x}_0, \boldsymbol{\epsilon}) = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon} \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

- Rewrite posterior mean

$$\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) \longrightarrow \tilde{\mu}_t\left(\mathbf{x}_t(\mathbf{x}_0, \boldsymbol{\epsilon}), \frac{1}{\sqrt{\bar{\alpha}_t}}(\mathbf{x}_t(\mathbf{x}_0, \boldsymbol{\epsilon}) - \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon})\right)$$

rewrote  $x_0$

$$\longrightarrow \frac{1}{\sqrt{\bar{\alpha}_t}} \left( \mathbf{x}_t(\mathbf{x}_0, \boldsymbol{\epsilon}) - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon} \right)$$

$$\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) := \frac{\sqrt{\bar{\alpha}_{t-1}} \beta_t}{1 - \bar{\alpha}_t} \mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{x}_t \quad \alpha_t := 1 - \beta_t \quad \bar{\alpha}_t := \prod_{s=1}^t \bar{\alpha}_s.$$

# Prediction model, sampling becomes

- Remember

$$\mu_\theta(x_t, t) = \tilde{\mu}_t(x_t, x_0)$$

our model

- Becomes now

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right)$$

- $\epsilon_\theta$  is available function approximator intended to predict  $\epsilon$  from  $\mathbf{x}_t$ : LEARN

- To sample  $\mathbf{x}_{t-1} \sim p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$  to compute:

$$\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}, \quad \mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

*phew*  
*(now algorithm ☺)*

# DPM training and sampling

---

## Algorithm 1 Training

---

```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
         $\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2$ 
6: until converged
```

---

predicted noise

$\mathbf{x}_t$ , the noised image at  
timestep  $t$

---

## Algorithm 2 Sampling

---

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  sample from prior
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \mathbf{z}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

---

predicted  
noise


add noise back for stochasticity  
(and ensure noise remains  
Gaussian)

# Example Diffusion Model *Generating Antibodies*

# Introduction to antibodies

- Antibody (**Ab**) or Immunoglobulin (**Ig**)  
Y-shaped protein produced by the immune system  
in response to **antigens (Ag)**


- Two **heavy (H)** and two **light (L)** chains
- **Constant regions** (mostly H)
- **Variable regions (H/L)**
  - Antigen-binding site (paratope)
  - Framework region + CDR loops
  - CDR-H3 loop** → most variable



- Monoclonal antibody (**mAb**) → engineered in the lab

# Diffusion model for antibody design – DiffAb

- In antibody design, diffusion models used to **generate CDR sequences and structures**
- Condition on **bound antibody-antigen complex**
- **Multimodal** → amino acid types,  $\text{Ca}$  atom coordinates, and orientations in  $\text{SO}(3)$




# DiffAb – Diffusion processes

Feature	Forward diffusion process ( $t = 0, \dots, T$ ). Distributions $q$	Prior distribution (sampling from $t = T$ )	Generate diffusion process ( $t = T, \dots, 0$ ). Neural network parameterization
Amino acid types	Multinomial distribution: $q(s_j^t   s_j^0) = \text{Multinomial}\left(\bar{\alpha}^t \cdot \text{onehot}(s_j^0) + \frac{1 - \bar{\alpha}^t}{20} \cdot \mathbf{1}\right)$	Uniform distribution over 20 classes	$F \rightarrow$ MLP decoder to predict the probabilities of the 20 amino acids
C $\alpha$ coordinates	Normal distribution: $q(\mathbf{x}_j^t   \mathbf{x}_j^0) = N\left(\mathbf{x}_j^t \mid \sqrt{\bar{\alpha}^t} \cdot \mathbf{x}_j^0, (1 - \bar{\alpha}^t)\mathbf{I}\right)$	Standard normal distribution	$G \rightarrow$ MLP decoder to predict the coordinate deviation wrt the current orientation in the local frame
Orientations in SO(3)	Isotropic Gaussian distribution in SO(3): $q(\mathbf{o}_j^t   \mathbf{o}_j^0) = IG_{SO(3)}\left(\mathbf{o}_j^t \mid \text{ScaleRot}\left(\sqrt{\bar{\alpha}^t} \cdot \mathbf{o}_j^0\right), 1 - \bar{\alpha}^t\right)$	Uniform distribution over SO(3)	$H \rightarrow$ MLP decoder to predict the so(3) vector that is converted to orientation matrix

- $F, G, H$  share an encoder (3D attention layers) of single and pairwise features from the previous timestep
- These networks are **equivariant to the rotation and translation of the overall structure**

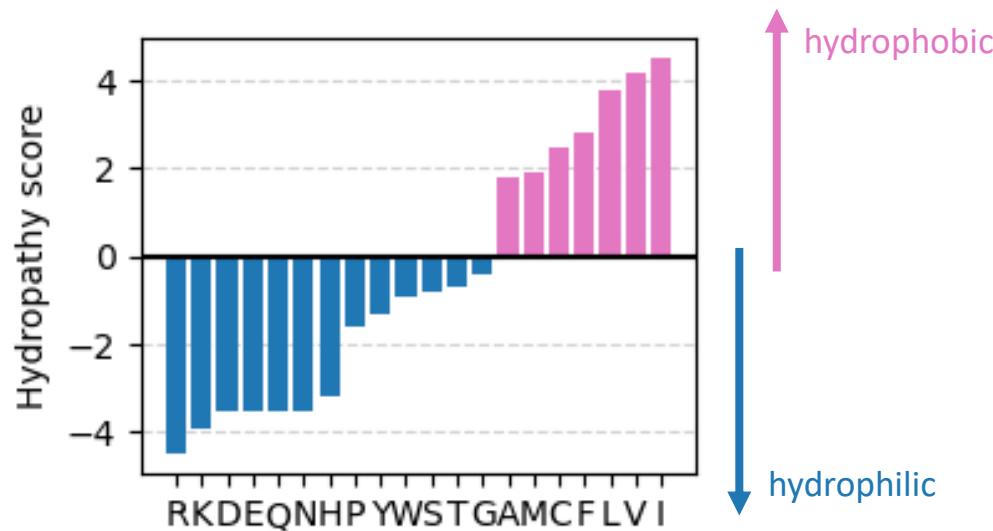
# Result: Generate new antibodies



# Developability properties


- **Which properties** are essential for antibody developability?
- How can we **calculate or predict** these properties?

## 1. Hydropathy score (proxy for solubility)



## 2. Folding energy ( $\Delta G$ )

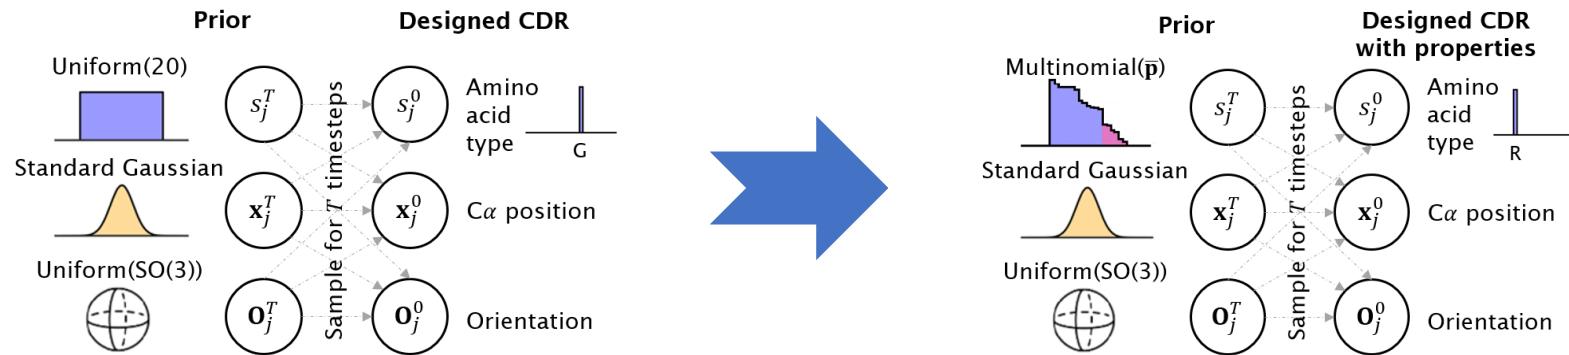
Predictor of changes in binding energy upon mutation ( $\Delta\Delta G$ ) for protein-protein complexes



J. Kyte and R.F. Doolittle. "A simple method for displaying the hydropathic character of a protein." *Journal of Molecular Biology*, 1982.

S. Shan, et al. "Deep learning guided optimization of human antibody against SARS-CoV-2 variants with broad neutralization." *Proceedings of the National Academy of Sciences*, 2022.


# Proposed guidance methods: Property-aware prior




- Sample from informative prior:

$$s_j^T \sim \text{Multinomial}(\bar{\mathbf{p}}) = (1 - b) \cdot \text{Uniform}(20) + b \cdot \text{Multinomial}(\mathbf{p})$$

- Probabilities  $\mathbf{p} = [p_1, \dots, p_{20}]$  for **hydropathy**




# Proposed guidance methods: Sampling by property



- At each **generation timestep**:
  1. Sample  $N$  times
  2. Select the sample with minimum property value
- With multiple properties, **select Pareto optimal** (minimum of the sum)
- Properties: **hydropathy score and predicted  $\Delta\Delta G$**

# Pareto optimal solutions

- Designs in the Pareto frontier
- Different CDR sequences lead to **similar structures** compared to the reference with **improved hydropathy and predicted  $\Delta\Delta G$**





END